



 Open in Colab

# AKUL VINOD

<https://colab.research.google.com/drive/17vhFouptDAaCYrUmMjlh1bbqGZWEItSj?usp=sharing>

## LoanTap Logistic Regression

**LoanTap** is an online platform committed to delivering customized loan products to millennials. They innovate in an otherwise dull loan segment, to deliver instant, flexible loans on consumer friendly terms to salaried professionals and businessmen.

The data science team at LoanTap is building an underwriting layer to determine the creditworthiness of MSMEs as well as individuals.

LoanTap deploys formal credit to salaried individuals and businesses 4 main financial instruments:

- 1 Personal Loan
- 2 EMI Free Loan
- 3 Personal Overdraft
- 4 Advance Salary Loan

This case study will focus on the underwriting process behind Personal Loan only

**Problem Statement:** Given a set of attributes for an individual, determine if a credit line should be extended to them. If so, what should the repayment terms be in business recommendations?

```
In [15]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
In [16]: pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
pd.set_option('display.width', 1000)
```

```
In [17]: df = pd.read_csv("logistic_regression.csv")
df.head()
```

```
Out[17]:
```

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	1
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	

```
In [18]: # Shape of the dataset
print("No. of rows : ", df.shape[0])
print("No. of columns : ", df.shape[1])
```

```
No. of rows : 396030
No. of columns : 27
```

```
In [19]: df.columns
```

```
Out[19]: Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade', 'sub_grade',
       'emp_title', 'emp_length', 'home_ownership', 'annual_inc', 'verification_status',
       'issue_d', 'loan_status', 'purpose', 'title', 'dti', 'earliest_cr_line',
       'open_acc', 'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'initial_list_status',
       'application_type', 'mort_acc', 'pub_rec_bankruptcies', 'addreses'],
      dtype='object')
```

```
In [20]: #No. of unique values for each columns
df.nunique()
```

Out[20]:

	<b>0</b>
<b>loan_amnt</b>	1397
<b>term</b>	2
<b>int_rate</b>	566
<b>installment</b>	55706
<b>grade</b>	7
<b>sub_grade</b>	35
<b>emp_title</b>	173105
<b>emp_length</b>	11
<b>home_ownership</b>	6
<b>annual_inc</b>	27197
<b>verification_status</b>	3
<b>issue_d</b>	115
<b>loan_status</b>	2
<b>purpose</b>	14
<b>title</b>	48816
<b>dti</b>	4262
<b>earliest_cr_line</b>	684
<b>open_acc</b>	61
<b>pub_rec</b>	20
<b>revol_bal</b>	55622
<b>revol_util</b>	1226
<b>total_acc</b>	118
<b>initial_list_status</b>	2
<b>application_type</b>	3
<b>mort_acc</b>	33
<b>pub_rec_bankruptcies</b>	9
<b>address</b>	393700

**dtype:** int64

In [21]: `df.isnull().sum()`

Out[21]:

	0
loan_amnt	0
term	0
int_rate	0
installment	0
grade	0
sub_grade	0
emp_title	22927
emp_length	18301
home_ownership	0
annual_inc	0
verification_status	0
issue_d	0
loan_status	0
purpose	0
title	1756
dti	0
earliest_cr_line	0
open_acc	0
pub_rec	0
revol_bal	0
revol_util	276
total_acc	0
initial_list_status	0
application_type	0
mort_acc	37795
pub_rec_bankruptcies	535
address	0

**dtype:** int64

In [22]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   loan_amnt        396030 non-null   float64
 1   term              396030 non-null   object 
 2   int_rate          396030 non-null   float64
 3   installment       396030 non-null   float64
 4   grade             396030 non-null   object 
 5   sub_grade          396030 non-null   object 
 6   emp_title          373103 non-null   object 
 7   emp_length         377729 non-null   object 
 8   home_ownership     396030 non-null   object 
 9   annual_inc         396030 non-null   float64
 10  verification_status 396030 non-null   object 
 11  issue_d            396030 non-null   object 
 12  loan_status         396030 non-null   object 
 13  purpose             396030 non-null   object 
 14  title               394274 non-null   object 
 15  dti                 396030 non-null   float64
 16  earliest_cr_line    396030 non-null   object 
 17  open_acc            396030 non-null   float64
 18  pub_rec              396030 non-null   float64
 19  revol_bal            396030 non-null   float64
 20  revol_util           395754 non-null   float64
 21  total_acc            396030 non-null   float64
 22  initial_list_status  396030 non-null   object 
 23  application_type      396030 non-null   object 
 24  mort_acc              358235 non-null   float64
 25  pub_rec_bankruptcies 395495 non-null   float64
 26  address              396030 non-null   object 
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
```

```
In [23]: #Converting string to date-time format
df['issue_d']=pd.to_datetime(df['issue_d'])
df['earliest_cr_line']=pd.to_datetime(df['earliest_cr_line'])
```

```
In [24]: #Need to fix the issue on title column
df['title'].value_counts()[:10]
```

Out[24]:

	count
	title
<b>Debt consolidation</b>	152472
<b>Credit card refinancing</b>	51487
<b>Home improvement</b>	15264
<b>Other</b>	12930
<b>Debt Consolidation</b>	11608
<b>Major purchase</b>	4769
<b>Consolidation</b>	3852
<b>debt consolidation</b>	3547
<b>Business</b>	2949
<b>Debt Consolidation Loan</b>	2864

**dtype:** int64

In [25]:

```
df['title']=df.title.str.lower()
df['title'].value_counts()[:10]
```

Out[25]:

	count
	title
<b>debt consolidation</b>	168108
<b>credit card refinancing</b>	51781
<b>home improvement</b>	17117
<b>other</b>	12993
<b>consolidation</b>	5583
<b>major purchase</b>	4998
<b>debt consolidation loan</b>	3513
<b>business</b>	3017
<b>medical expenses</b>	2820
<b>credit card consolidation</b>	2638

**dtype:** int64

In [26]:

```
# Extraction of pincode from the 'Address' column
df['pin_code'] = df['address'].str.split(' ').str[-1]
```

```
# Let's fetch ZIP from address and then drop the remaining details -  
df['zip_code'] = df['address'].fillna(' ').apply(lambda x: x[-5:])
```

```
In [27]: #Convert columns with less number of unique values to categorical columns  
cat_cols = ['term', 'grade','sub_grade','home_ownership',  
            'verification_status','loan_status','purpose',  
            'initial_list_status','application_type']  
  
df[cat_cols] = df[cat_cols].astype('category')
```

```
In [28]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 396030 entries, 0 to 396029  
Data columns (total 29 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   loan_amnt        396030 non-null  float64  
 1   term             396030 non-null  category  
 2   int_rate          396030 non-null  float64  
 3   installment       396030 non-null  float64  
 4   grade             396030 non-null  category  
 5   sub_grade          396030 non-null  category  
 6   emp_title          373103 non-null  object  
 7   emp_length         377729 non-null  object  
 8   home_ownership     396030 non-null  category  
 9   annual_inc         396030 non-null  float64  
 10  verification_status 396030 non-null  category  
 11  issue_d            396030 non-null  datetime64[ns]  
 12  loan_status         396030 non-null  category  
 13  purpose             396030 non-null  category  
 14  title               394274 non-null  object  
 15  dti                 396030 non-null  float64  
 16  earliest_cr_line    396030 non-null  datetime64[ns]  
 17  open_acc             396030 non-null  float64  
 18  pub_rec              396030 non-null  float64  
 19  revol_bal            396030 non-null  float64  
 20  revol_util            395754 non-null  float64  
 21  total_acc             396030 non-null  float64  
 22  initial_list_status   396030 non-null  category  
 23  application_type      396030 non-null  category  
 24  mort_acc              358235 non-null  float64  
 25  pub_rec_bankruptcies  395495 non-null  float64  
 26  address               396030 non-null  object  
 27  pin_code              396030 non-null  object  
 28  zip_code               396030 non-null  object  
dtypes: category(9), datetime64[ns](2), float64(12), object(6)  
memory usage: 63.8+ MB
```

```
In [29]: # Statistical summary of the dataset  
df.describe().round(2)
```

Out[29]:

	<b>loan_amnt</b>	<b>int_rate</b>	<b>installment</b>	<b>annual_inc</b>	<b>issue_d</b>	
<b>count</b>	396030.00	396030.00	396030.00	396030.00	396030	396030
<b>mean</b>	14113.89	13.64	431.85	74203.18	2014-02-02 15:57:58.045602560	1
<b>min</b>	500.00	5.32	16.08	0.00	2007-06-01 00:00:00	1
<b>25%</b>	8000.00	10.49	250.33	45000.00	2013-05-01 00:00:00	1
<b>50%</b>	12000.00	13.33	375.43	64000.00	2014-04-01 00:00:00	1
<b>75%</b>	20000.00	16.49	567.30	90000.00	2015-03-01 00:00:00	2
<b>max</b>	40000.00	30.99	1533.81	8706582.00	2016-12-01 00:00:00	9999
<b>std</b>	8357.44	4.47	250.73	61637.62	NaN	18

- Nearly 80% of the loans have a term of 36 months.
- The majority of loans (30%) are graded as B, followed by C, A, and D respectively.
- For 50% of cases, the type of home ownership is mortgage.
- The loan status target variable is biased towards fully-paid loans, with defaulters accounting for approximately 25% of fully-paid instances.
- Approximately 85% of applicants do not have a public record or have not filed for bankruptcy.
- Nearly all applicants (99%) have applied under the 'individual' application type. -The most common purpose for taking out loans is debt consolidation, accounting for 55%, followed by 20% for credit card purposes.

In [30]: `# Checking the distribution of the outcome labels  
df.loan_status.value_counts(normalize=True)*100`

Out[30]:

	<b>proportion</b>
<b>loan_status</b>	
<b>Fully Paid</b>	80.387092
<b>Charged Off</b>	19.612908

**dtype:** float64

- 80% belongs to the class 0 : which is loan fully paid.
- 20% belongs to the class 1 : which were charged off.

**As we can see, there is an imbalance in the data**

```
In [31]: df.initial_list_status.value_counts(normalize=True)*100
```

```
Out[31]: proportion
```

initial_list_status	
f	60.113123
w	39.886877

**dtype:** float64

- 60% belongs to whole loans(w)
- 40% belongs to fractional loans(f)

```
In [32]: df.application_type.value_counts(normalize=True)*100
```

```
Out[32]: proportion
```

application_type	
INDIVIDUAL	99.820468
JOINT	0.107315
DIRECT_PAY	0.072217

**dtype:** float64

- Maximum belongs to INDIVIDUAL application

```
In [33]: df.term.value_counts(normalize=True)*100
```

```
Out[33]: proportion
```

term	
36 months	76.258112
60 months	23.741888

**dtype:** float64

36-month loan terms apply to 76% of the loans, while 60-month loan terms apply to 24% of them.

```
In [34]: # The home ownership status provided by the borrower during registration or ob  
df.home_ownership.value_counts(normalize=True)*100
```

Out[34]: **proportion**

home_ownership	proportion
<b>MORTGAGE</b>	50.084085
<b>RENT</b>	40.347953
<b>OWN</b>	9.531096
<b>OTHER</b>	0.028281
<b>NONE</b>	0.007828
<b>ANY</b>	0.000758

**dtype:** float64

- 50% loans belongs to mortgage
- 40% loans belongs to rent

```
In [35]: df.verification_status.value_counts(normalize=True)*100
```

Out[35]: **proportion**

verification_status	proportion
<b>Verified</b>	35.240512
<b>Source Verified</b>	33.175517
<b>Not Verified</b>	31.583971

**dtype:** float64

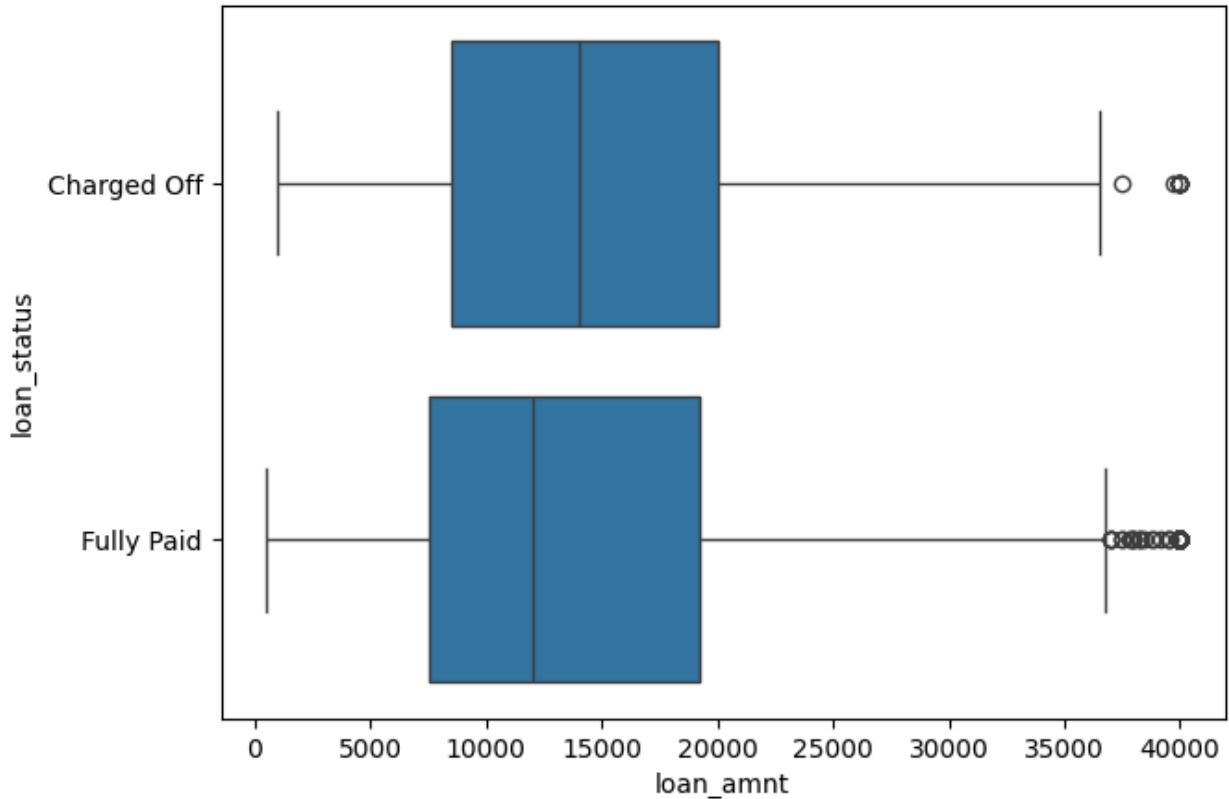
```
In [36]: df.groupby(by = "loan_status")["loan_amnt"].describe()
```

Out[36]:

loan_status	count	mean	std	min	25%	50%	75%
<b>Charged Off</b>	77673.0	15126.300967	8505.090557	1000.0	8525.0	14000.0	20000.
<b>Fully Paid</b>	318357.0	13866.878771	8302.319699	500.0	7500.0	12000.0	19225.

```
In [37]: plt.figure(figsize=(7,5))
sns.boxplot(x=df["loan_amnt"], y=df["loan_status"])
```

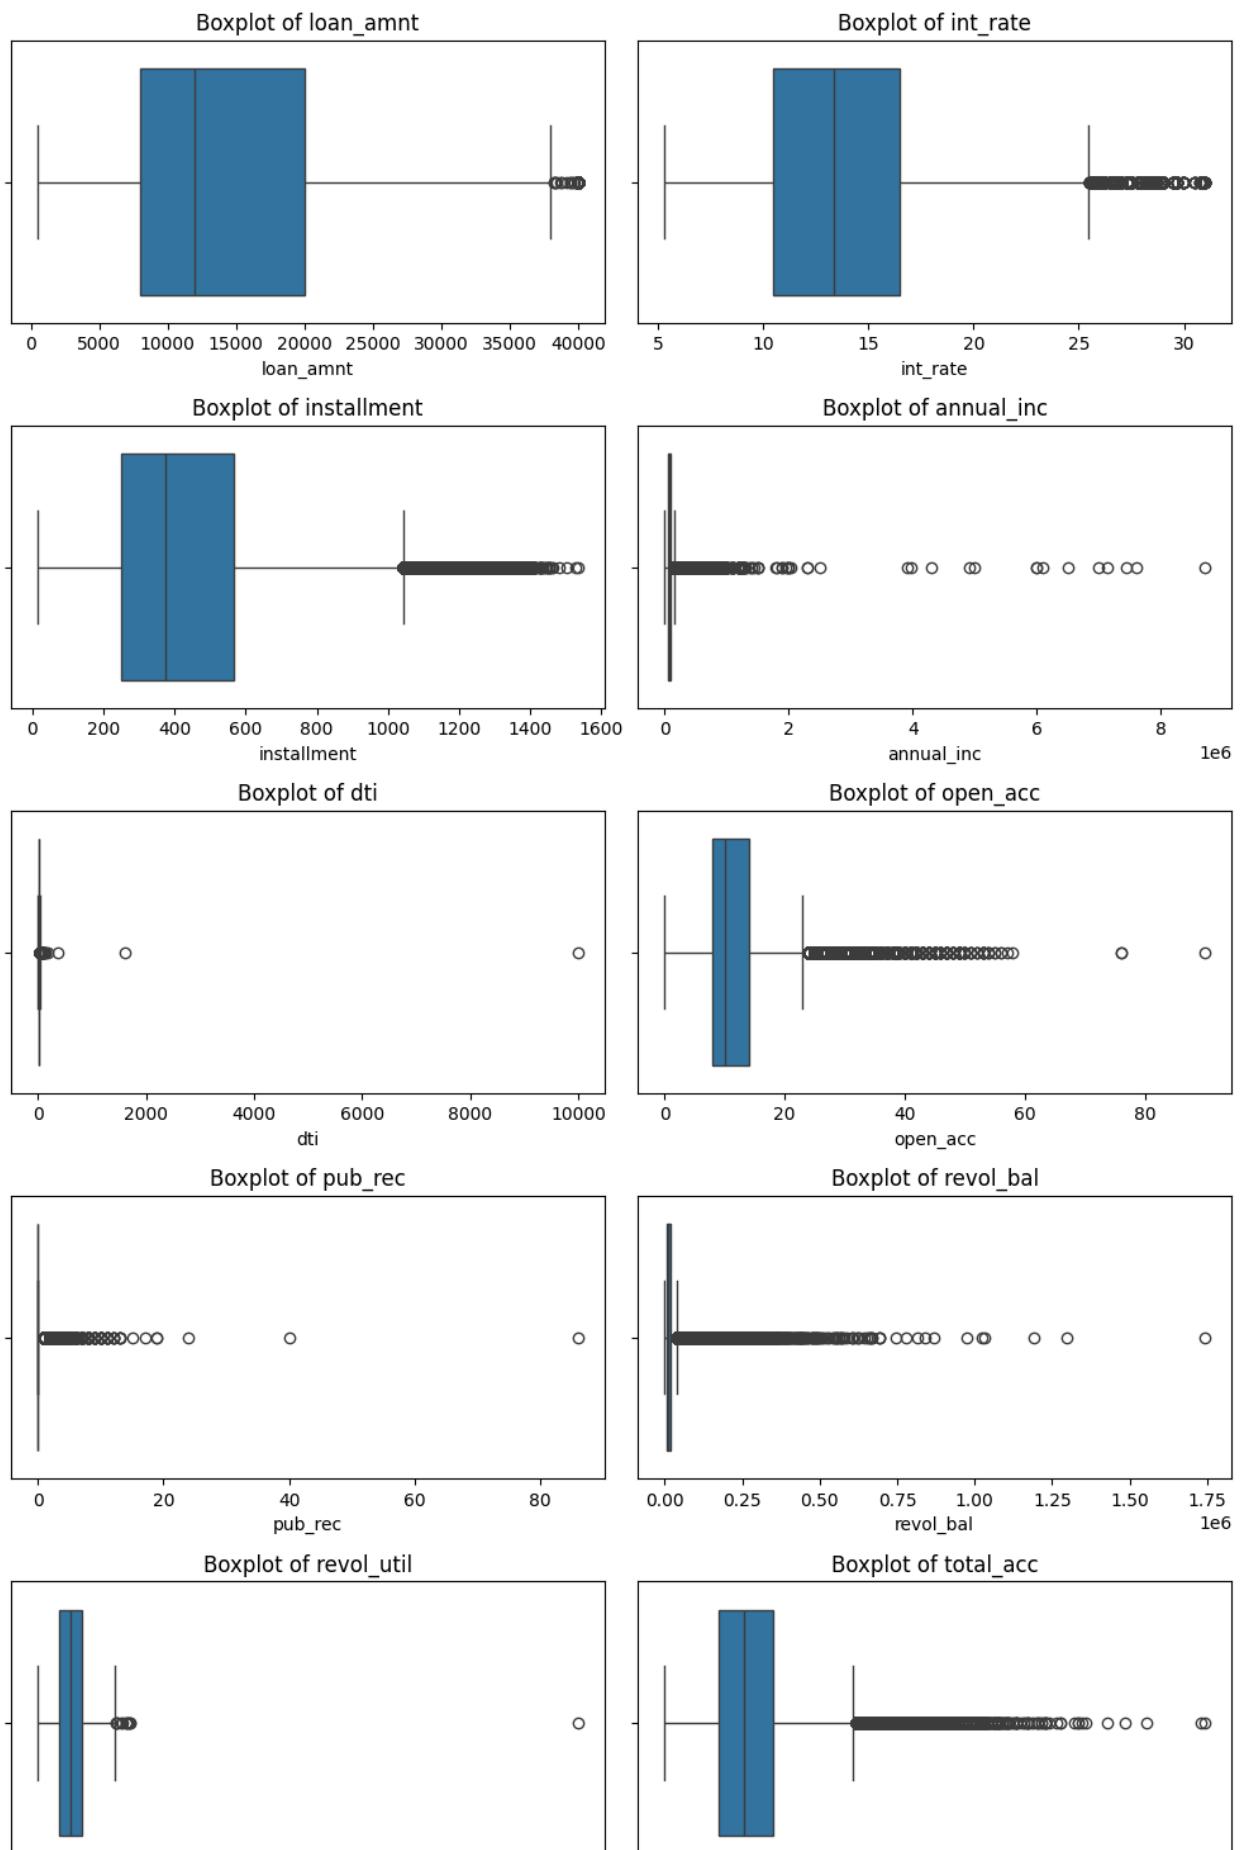
```
Out[37]: <Axes: xlabel='loan_amnt', ylabel='loan_status'>
```



```
In [38]: num_cols = df.select_dtypes(include='number').columns

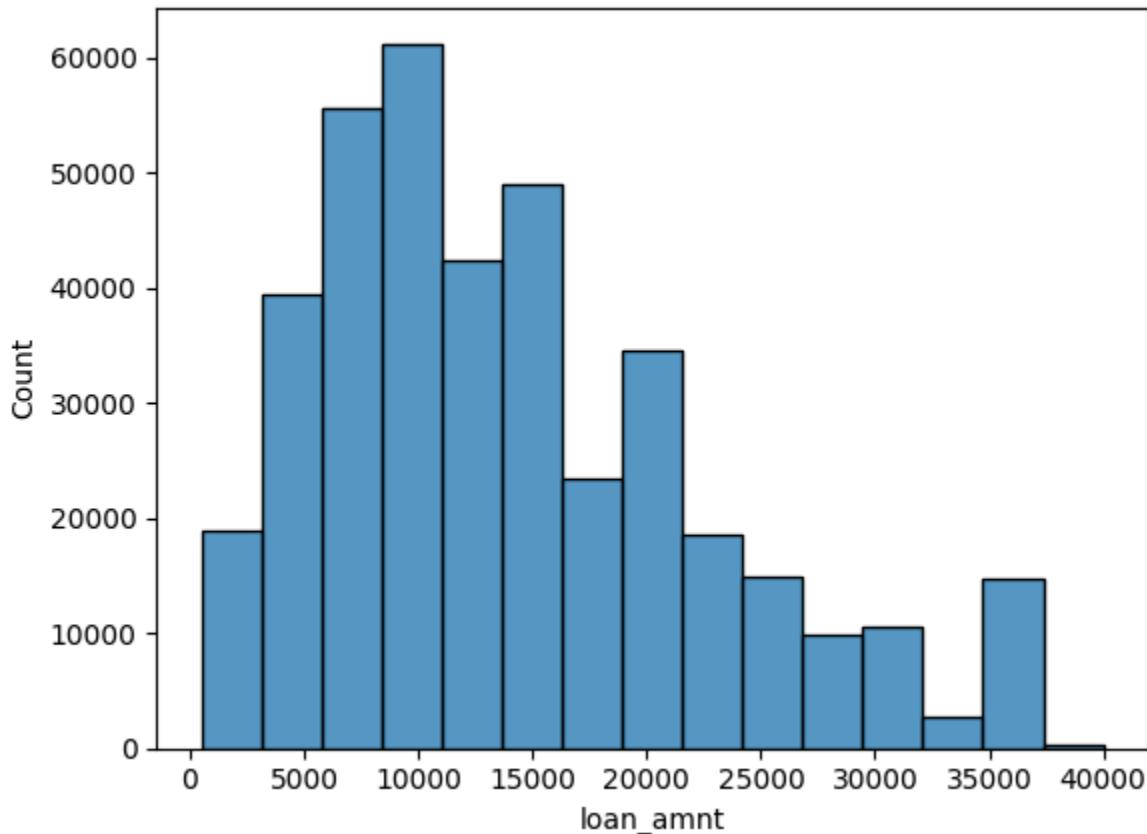
fig = plt.figure(figsize=(10,21))
i=1
for col in num_cols:
    ax = plt.subplot(7,2,i)
    sns.boxplot(x=df[col])
    plt.title(f'Boxplot of {col}')
    i += 1

plt.tight_layout()
plt.show()
```



```
In [39]: sns.histplot(df["loan_amnt"], bins = 15)
```

```
Out[39]: <Axes: xlabel='loan_amnt', ylabel='Count'>
```



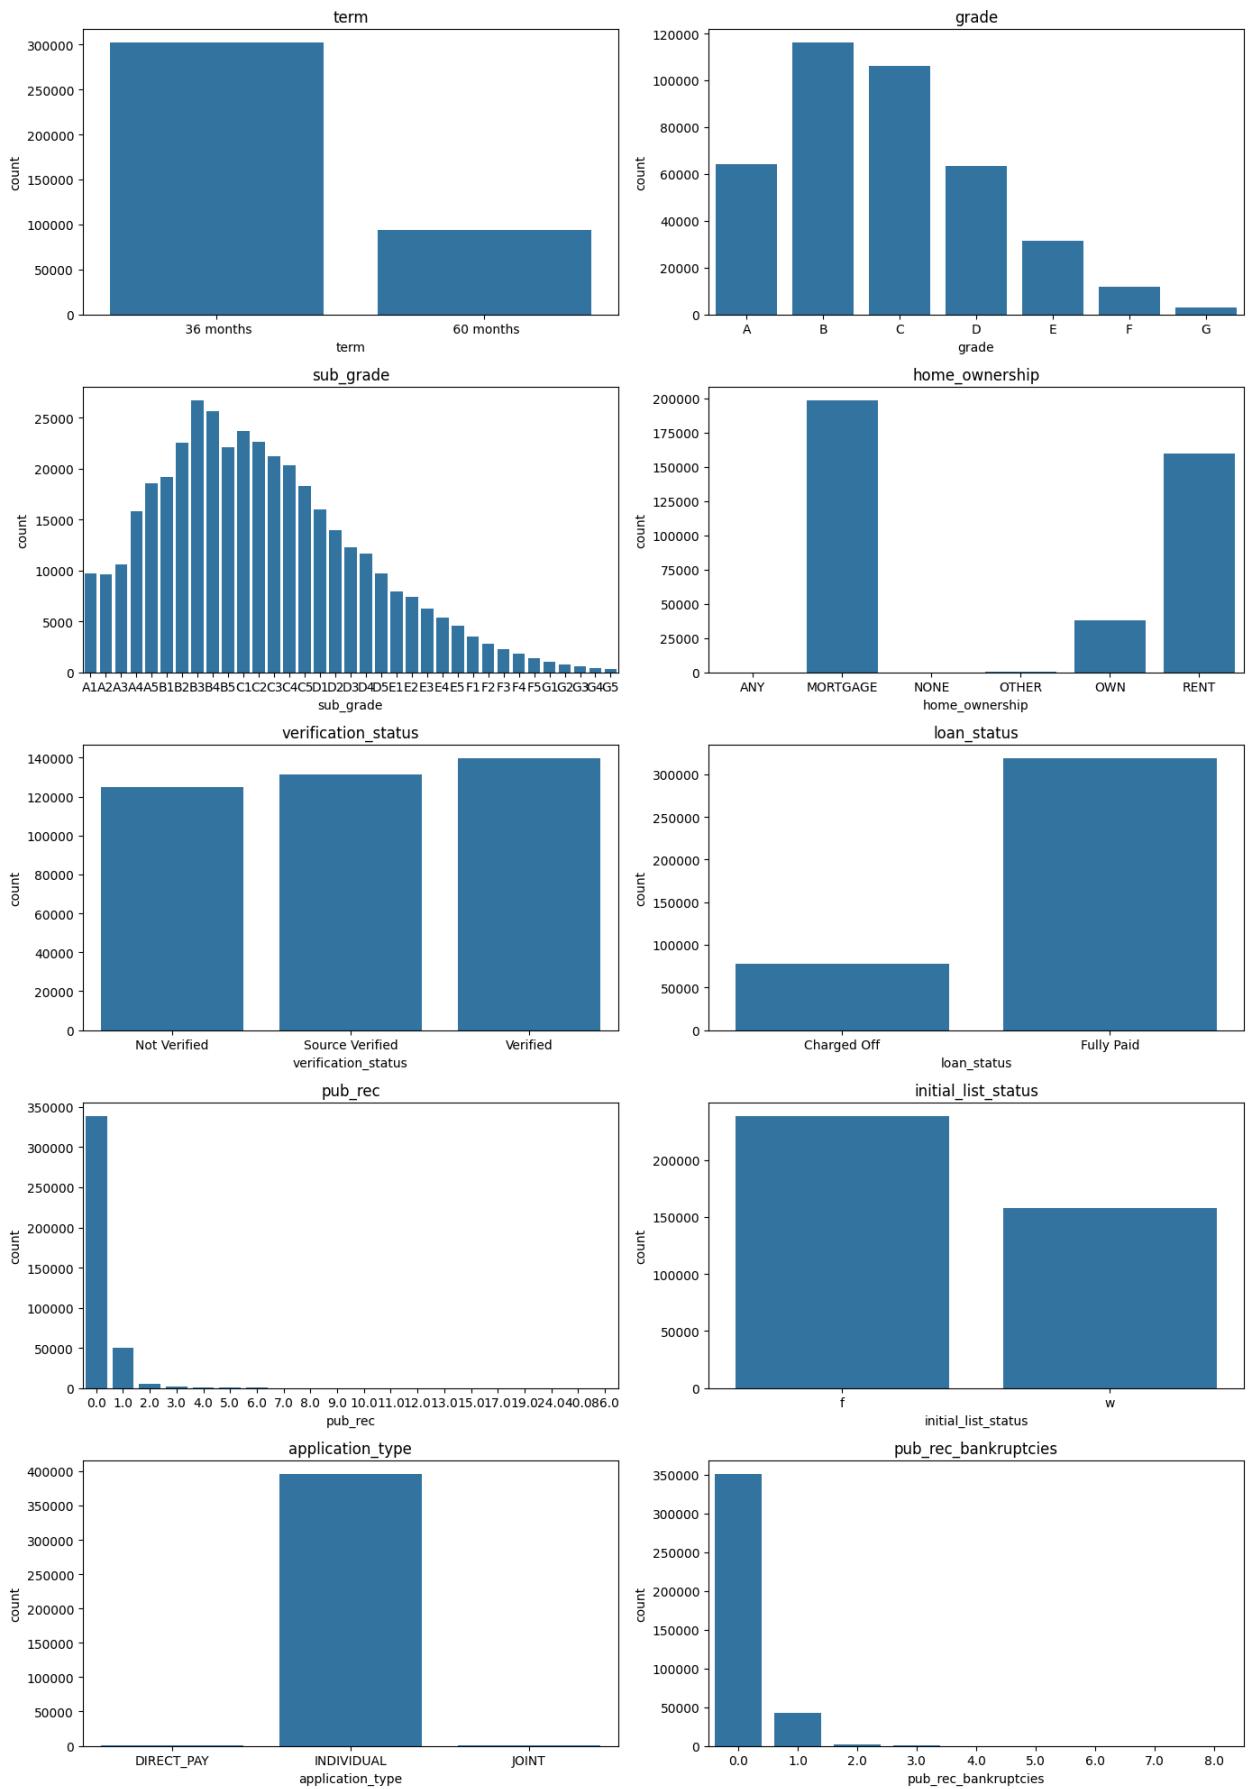
- for loan status Charged\_off, the mean and median of loan\_amount is higher than fully paid.
- also the distribution of loan\_amnt is right skewed, which says it has outlier presence.

```
In [40]: #Distribution of categorical variables
```

```
plot = ['term', 'grade', 'sub_grade', 'home_ownership', 'verification_status',
        'loan_status', 'pub_rec', 'initial_list_status',
        'application_type', 'pub_rec_bankruptcies']

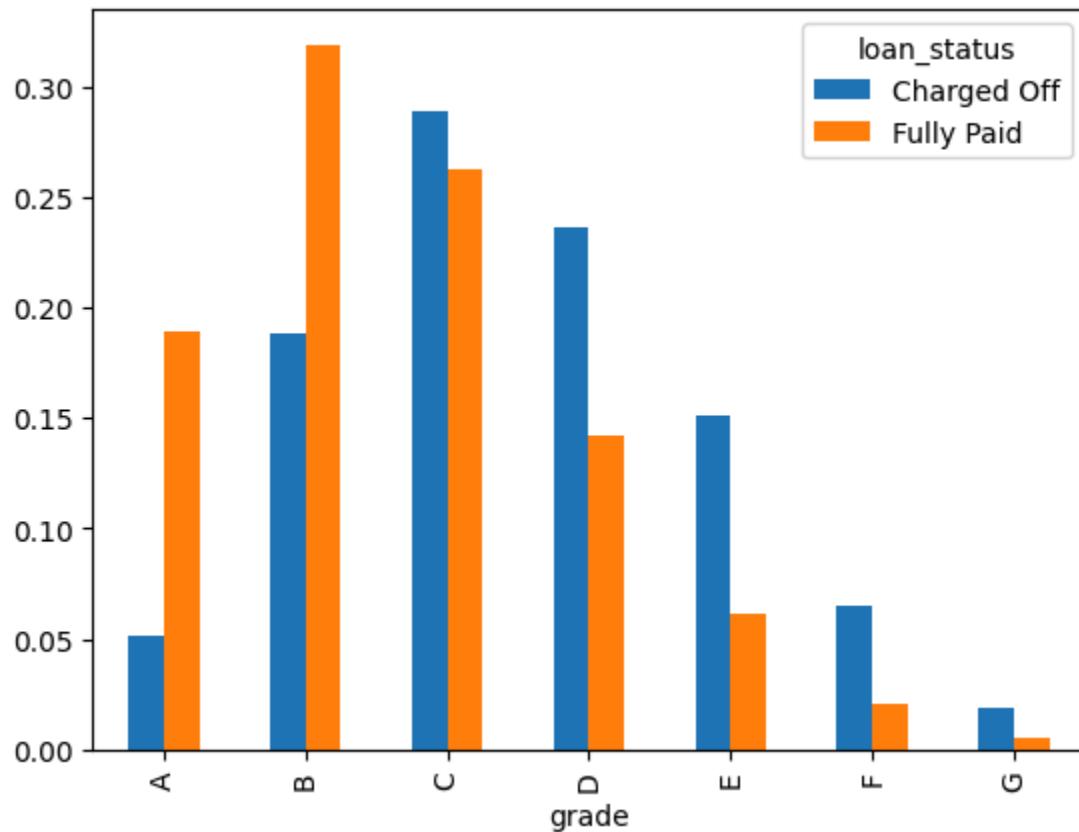
plt.figure(figsize=(14,20))
i=1
for col in plot:
    ax=plt.subplot(5,2,i)
    sns.countplot(x=df[col])
    plt.title(f'{col}')
    i += 1

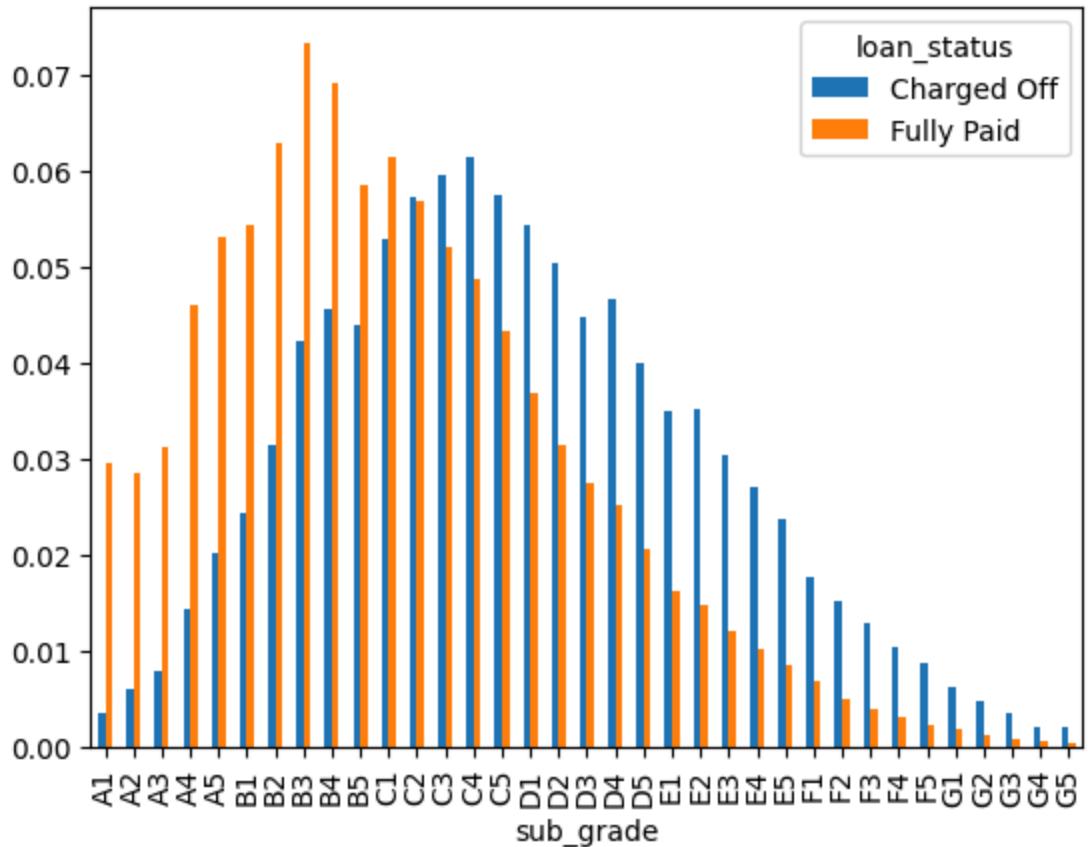
plt.tight_layout()
plt.show()
```



```
In [41]: pd.crosstab(index = df["grade"], columns= df["loan_status"],normalize= "columns")  
pd.crosstab(index = df["sub_grade"], columns= df["loan_status"],normalize= "columns")
```

```
Out[41]: <Axes: xlabel='sub_grade'>
```





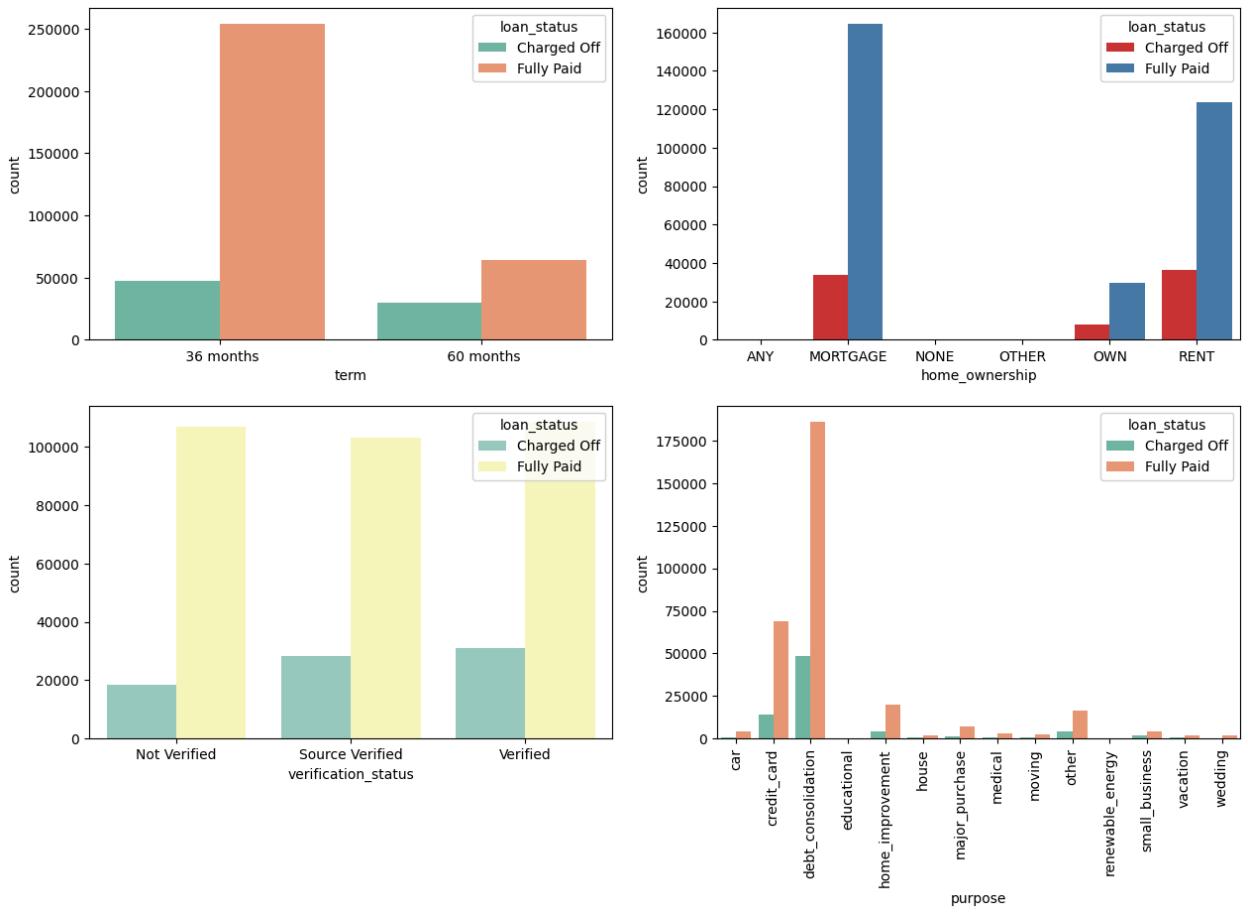
```
In [42]: plt.figure(figsize=(15,20))

plt.subplot(4,2,1)
sns.countplot(x='term', data=df, hue='loan_status', palette='Set2')

plt.subplot(4,2,2)
sns.countplot(x='home_ownership', data=df, hue='loan_status', palette='Set1')

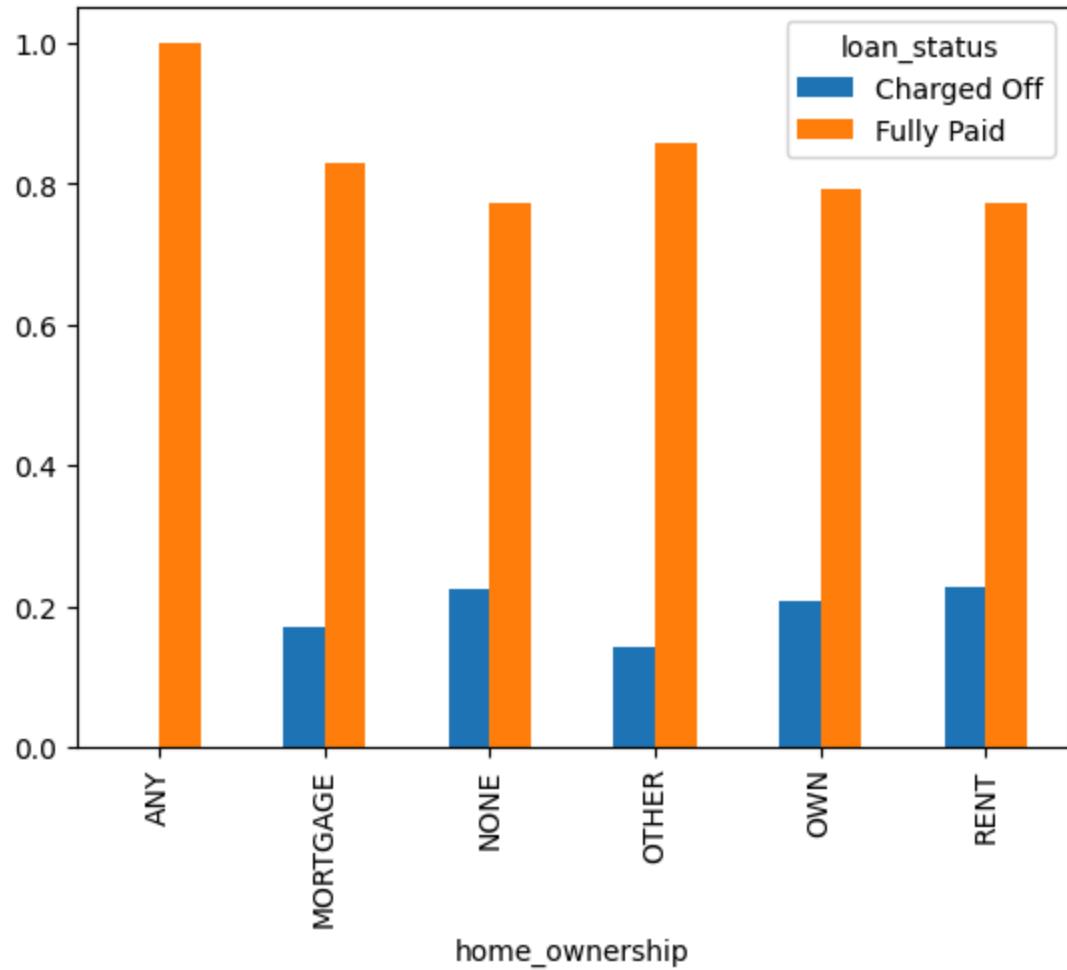
plt.subplot(4,2,3)
sns.countplot(x='verification_status', data=df, hue='loan_status', palette='Set3')

plt.subplot(4,2,4)
g=sns.countplot(x='purpose', data=df, hue='loan_status', palette='Set2')
g.set_xticklabels(g.get_xticklabels(), rotation=90)
plt.show()
```



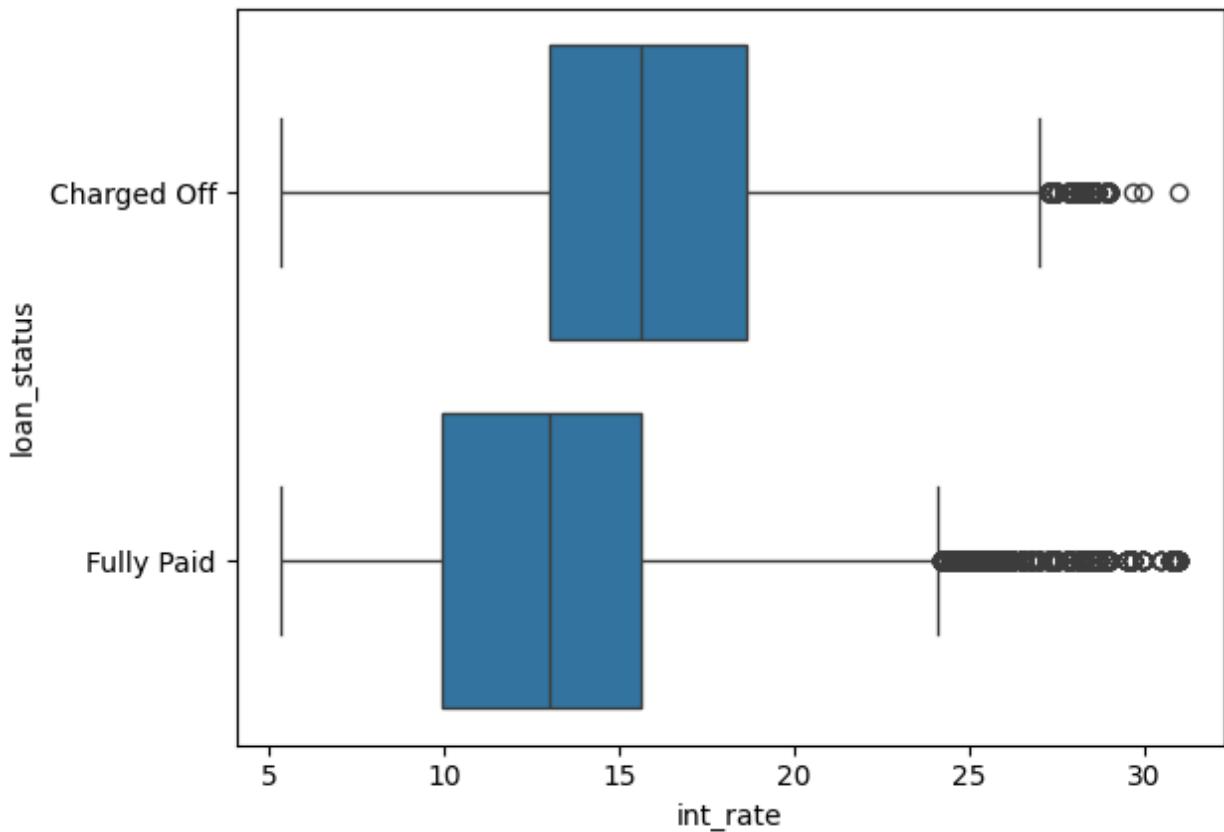
```
In [43]: pd.crosstab(index = df["home_ownership"], columns= df["loan_status"], normalize=True)
```

```
Out[43]: <Axes: xlabel='home_ownership'>
```

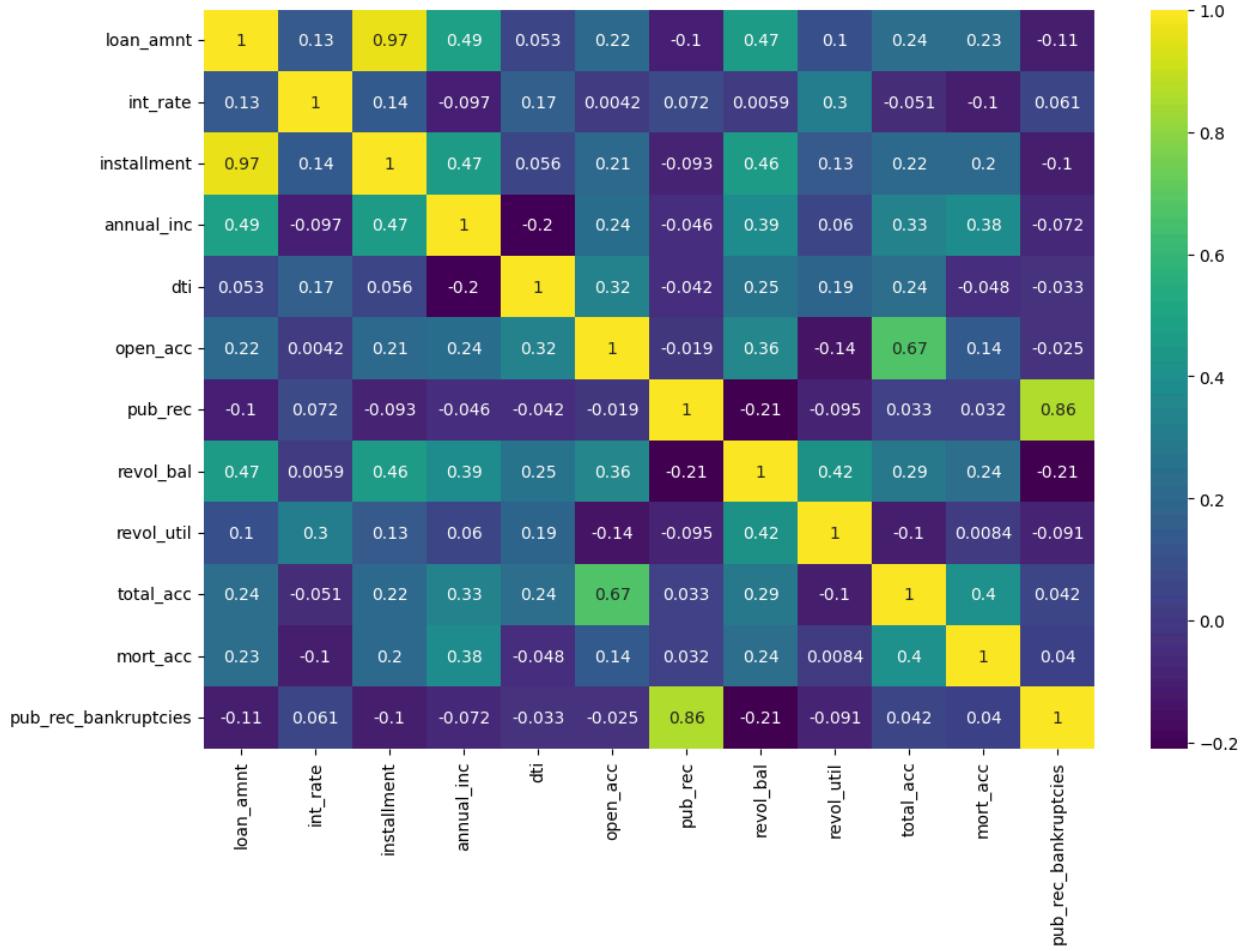


```
In [44]: sns.boxplot(x=df["int_rate"],y=df["loan_status"])
```

```
Out[44]: <Axes: xlabel='int_rate', ylabel='loan_status'>
```



```
In [46]: numerical_df = df.select_dtypes(include=np.number)
plt.figure(figsize=(12, 8))
sns.heatmap(numerical_df.corr(method='spearman'), annot=True, cmap='viridis')
plt.show()
```



## Data Processing

- In both the 'Categorical' and 'Numerical' categories, there are columns with a significant amount of missing data.
- For the 'Numerical' data, these missing values will be filled with the mean, while for the 'Categorical' data, they will be filled with the mode.

```
In [47]: #Check for Duplicate Values
df.duplicated().sum()
```

```
Out[47]: np.int64(0)
```

```
In [48]: # Null values replaced by 'Mode' in case of 'Categorical' column.
column_mode = df['emp_length'].mode()[0]
df['emp_length'] = df['emp_length'].fillna(column_mode)

# Null values replaced by 'Mean' in case of 'Numerical' column.
for column in ['revol_util', 'mort_acc', 'pub_rec_bankruptcies']:
    column_mean = df[column].mean()
    df[column] = df[column].fillna(column_mean)
```

```
In [49]: df.isna().sum()
```

Out[49]:

	<b>0</b>
loan_amnt	0
term	0
int_rate	0
installment	0
grade	0
sub_grade	0
emp_title	22927
emp_length	0
home_ownership	0
annual_inc	0
verification_status	0
issue_d	0
loan_status	0
purpose	0
title	1756
dti	0
earliest_cr_line	0
open_acc	0
pub_rec	0
revol_bal	0
revol_util	0
total_acc	0
initial_list_status	0
application_type	0
mort_acc	0
pub_rec_bankruptcies	0
address	0
pin_code	0
zip_code	0

**dtype:** int64

```
In [50]: # Dropping some variables which we can let go for now
df.drop(columns=['issue_d', 'emp_title', 'title', 'sub_grade',
                 'address', 'earliest_cr_line', 'emp_length'],
        axis=1, inplace=True)
```

## Feature Engineering

```
In [51]: def pub_rec(number):
    if number == 0:
        return 0
    else:
        return 1

def mort_acc(number):
    if number == 0.0:
        return 0
    elif number >= 1.0:
        return 1
    else:
        return number

def pub_rec_bankruptcies(number):
    if number == 0.0:
        return 0
    elif number >= 1.0:
        return 1
    else:
        return number
```

```
In [52]: df['pub_rec']=df.pub_rec.apply(pub_rec)
df['mort_acc']=df.mort_acc.apply(mort_acc)
df['pub_rec_bankruptcies']=df.pub_rec_bankruptcies.apply(pub_rec_bankruptcies)
```

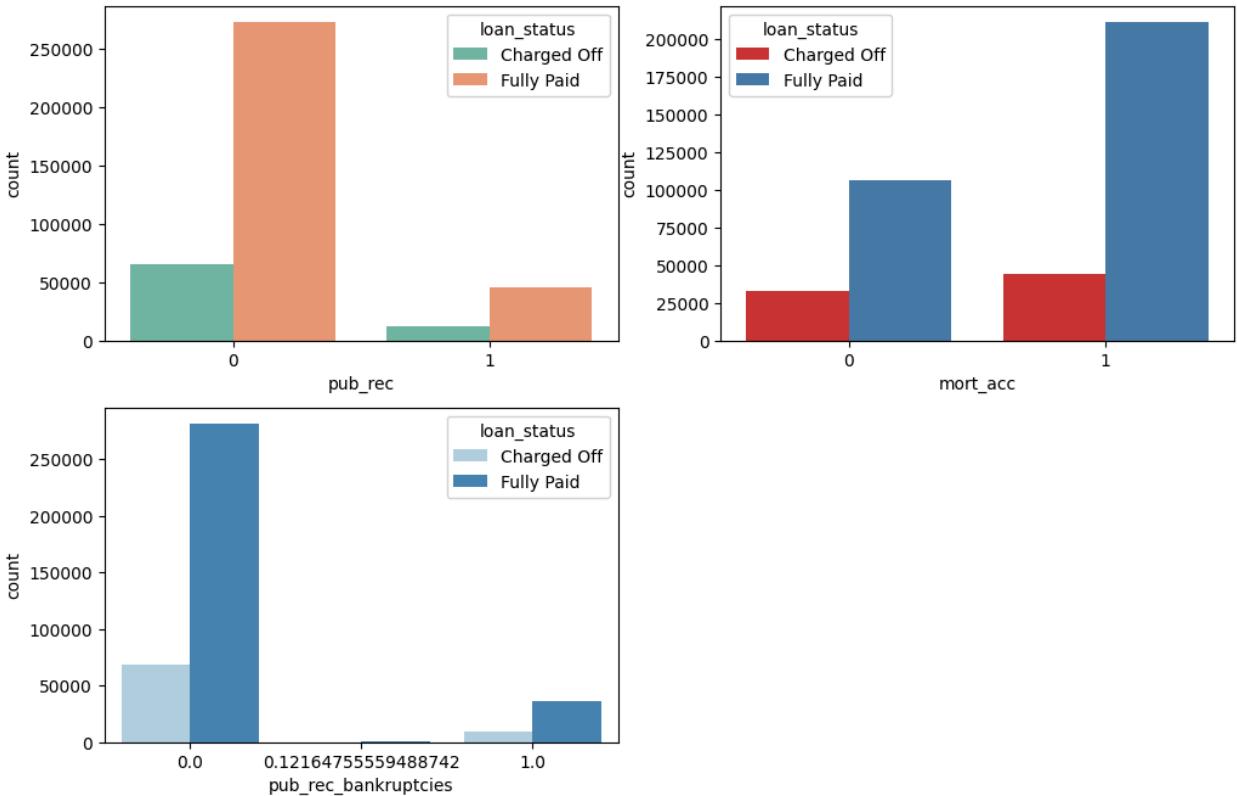
```
In [53]: plt.figure(figsize=(12,25))

plt.subplot(6,2,1)
sns.countplot(x='pub_rec',data=df,hue='loan_status', palette='Set2')

plt.subplot(6,2,2)
sns.countplot(x='mort_acc',data=df,hue='loan_status', palette='Set1')

plt.subplot(6,2,3)
sns.countplot(x='pub_rec_bankruptcies',data=df,hue='loan_status', palette='Blu')
```

```
Out[53]: <Axes: xlabel='pub_rec_bankruptcies', ylabel='count'>
```



```
In [54]: #pre processing
df['loan_status']=df.loan_status.map({'Fully Paid':0, 'Charged Off':1})

term_values={' 36 months': 36, ' 60 months':60}
df['term'] = df.term.map(term_values)

list_status = {'w': 0, 'f': 1}
df['initial_list_status'] = df.initial_list_status.map(list_status)
```

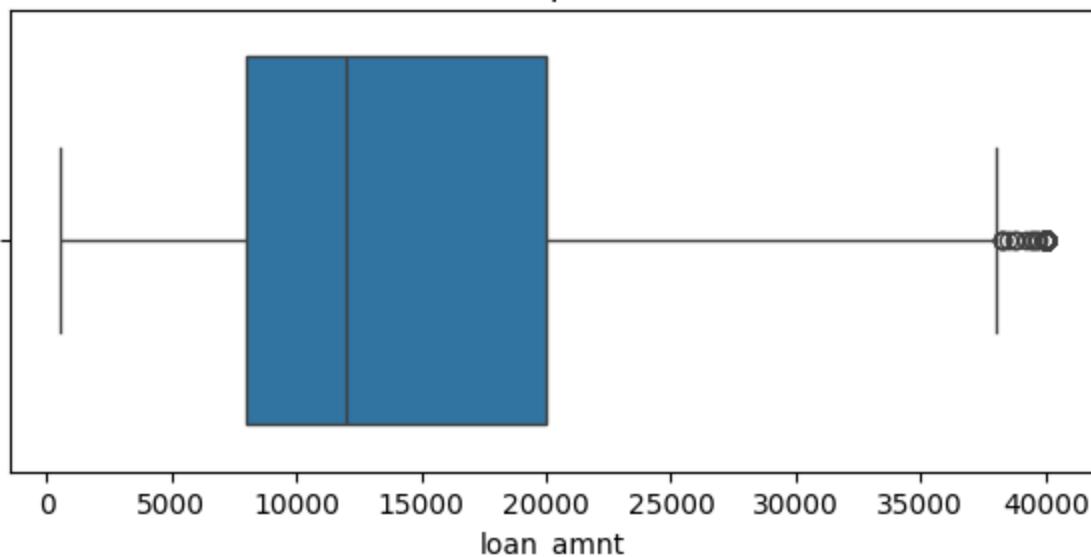
## Outlier Detection & Treatment

```
In [55]: numerical_col=df.select_dtypes(include='number')
num_cols=numerical_col.columns

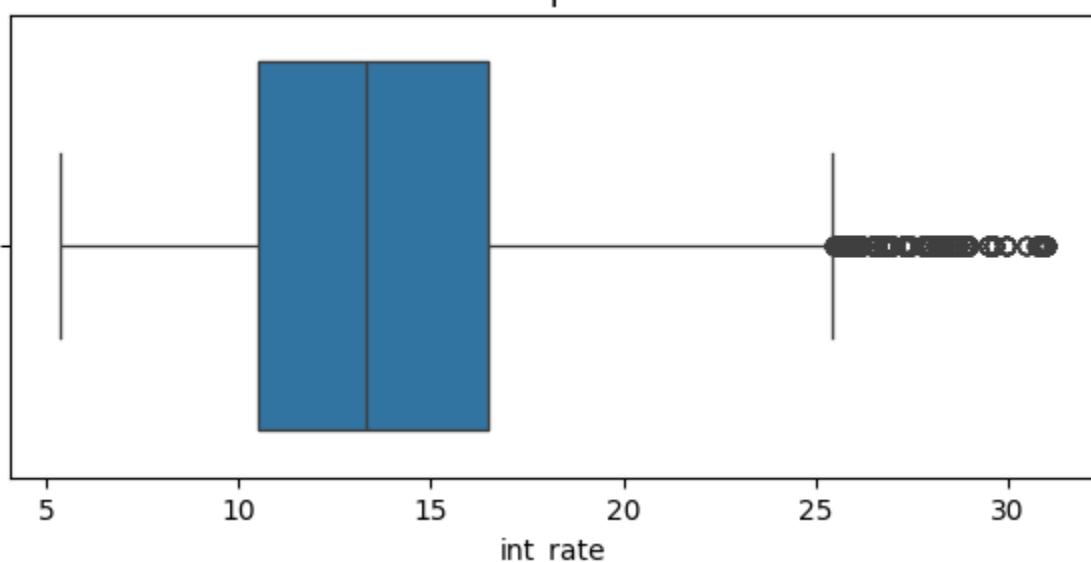
def box_plot(col):
    plt.figure(figsize=(7,3))
    sns.boxplot(x=df[col])
    plt.title('Boxplot')
    plt.show()

for col in num_cols:
    box_plot(col)
```

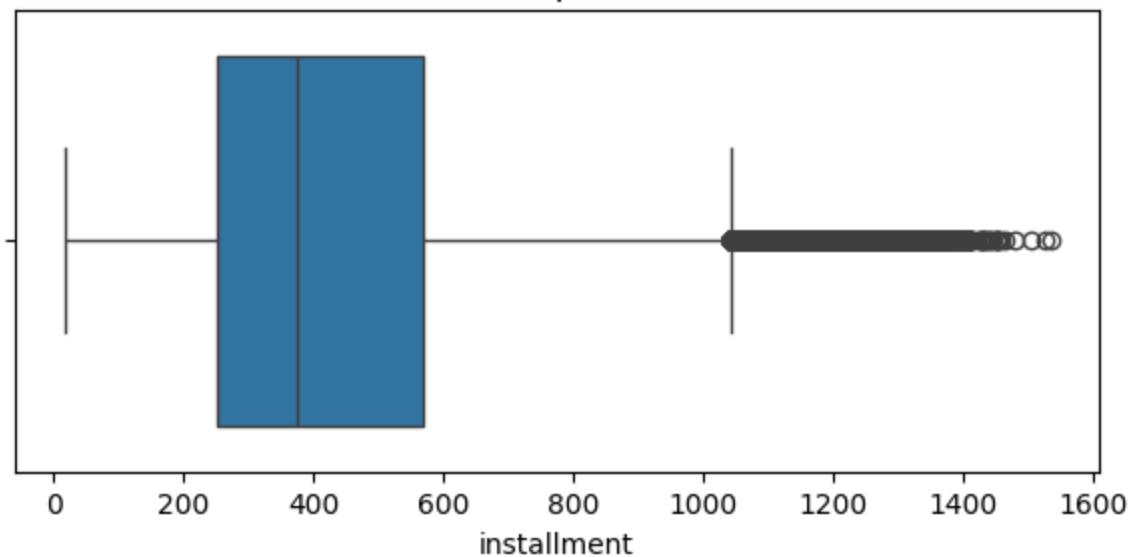
Boxplot



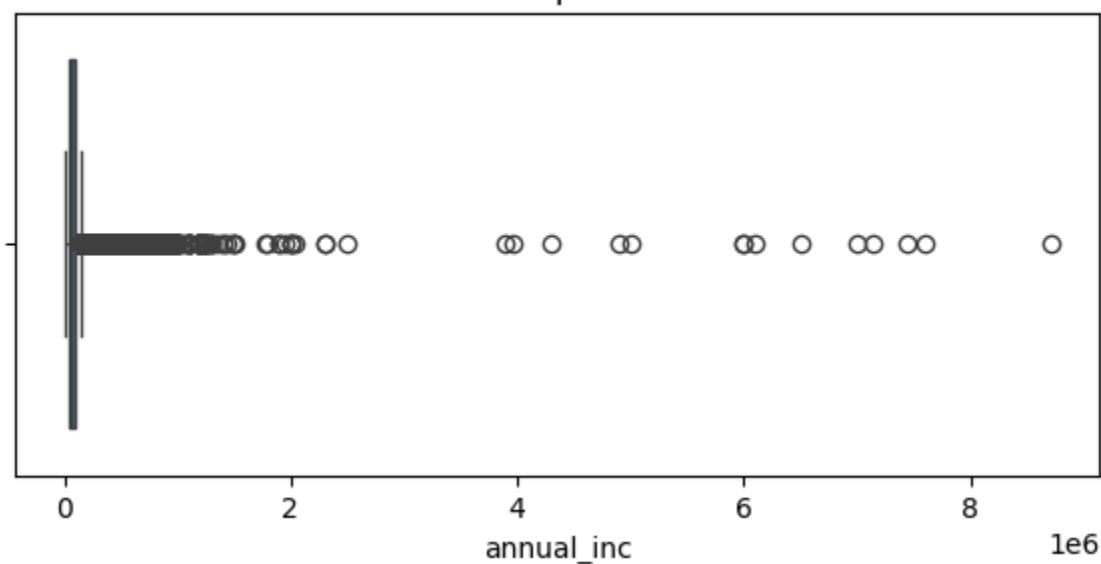
Boxplot



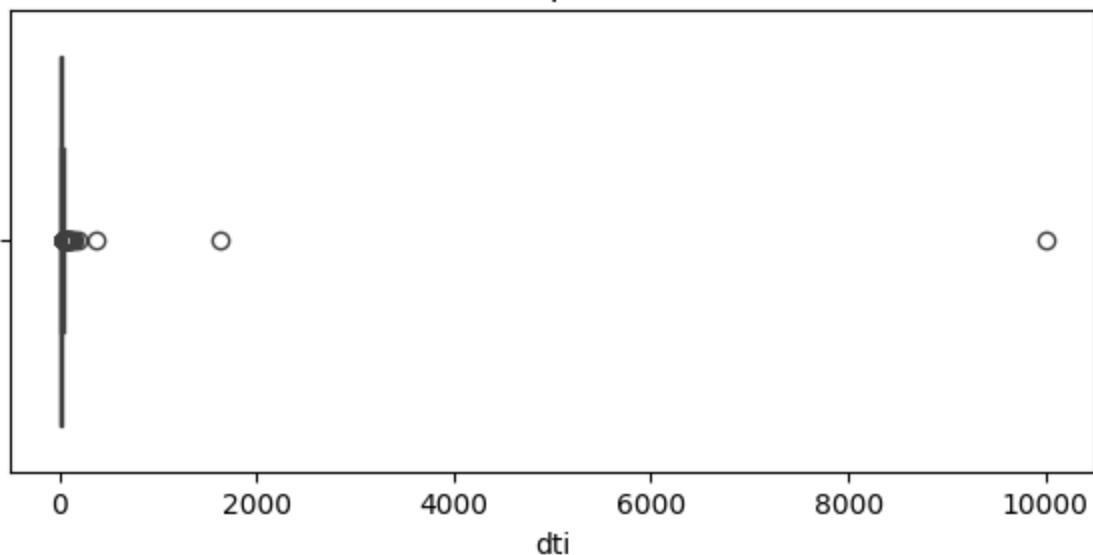
Boxplot



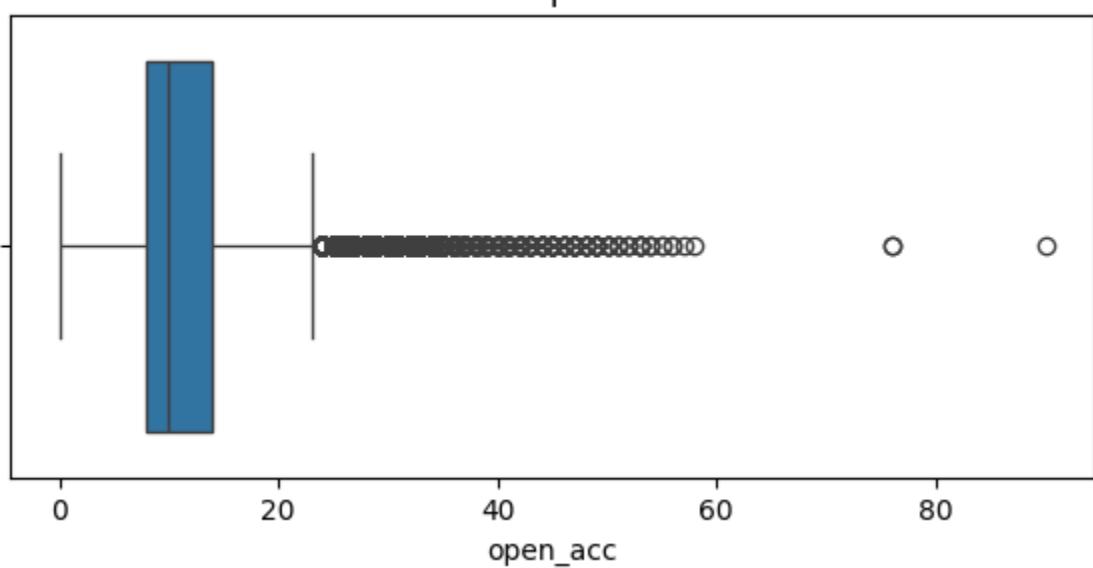
Boxplot



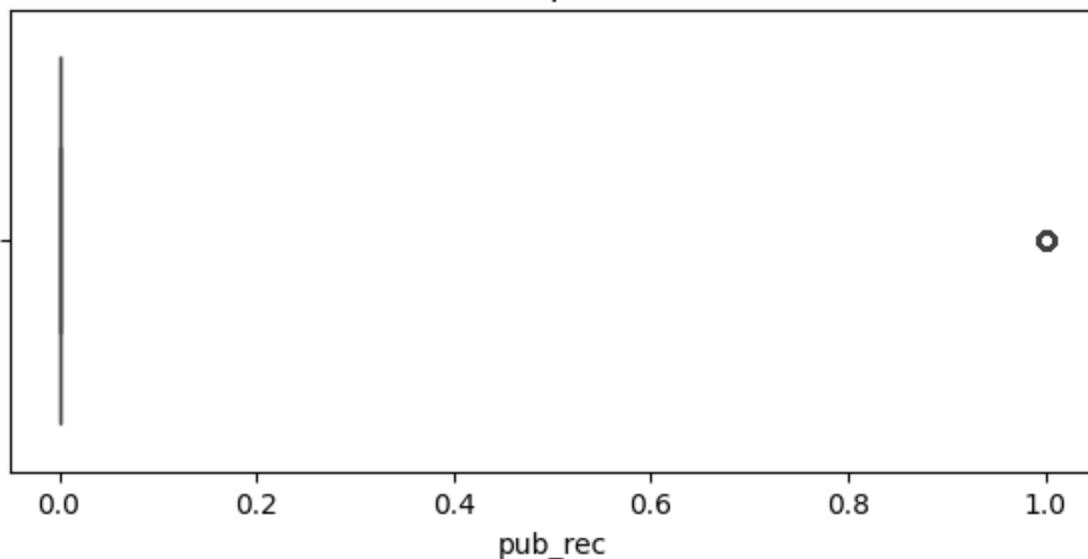
Boxplot



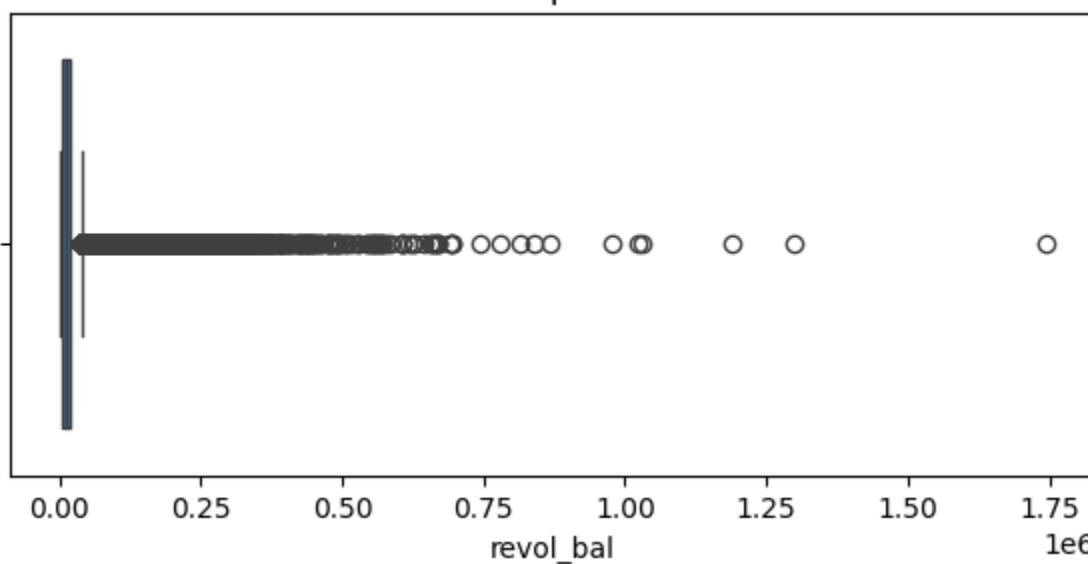
Boxplot



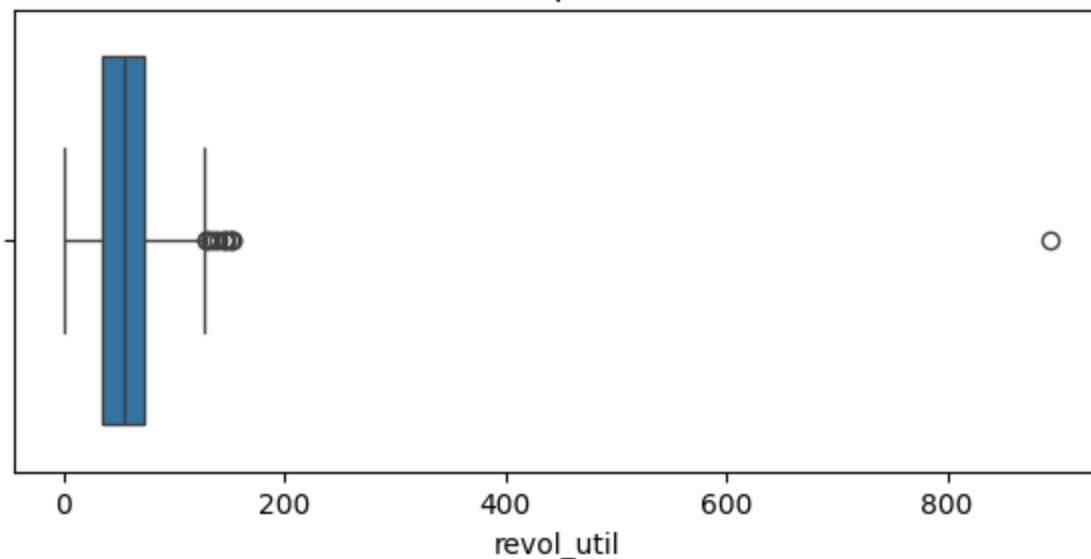
Boxplot



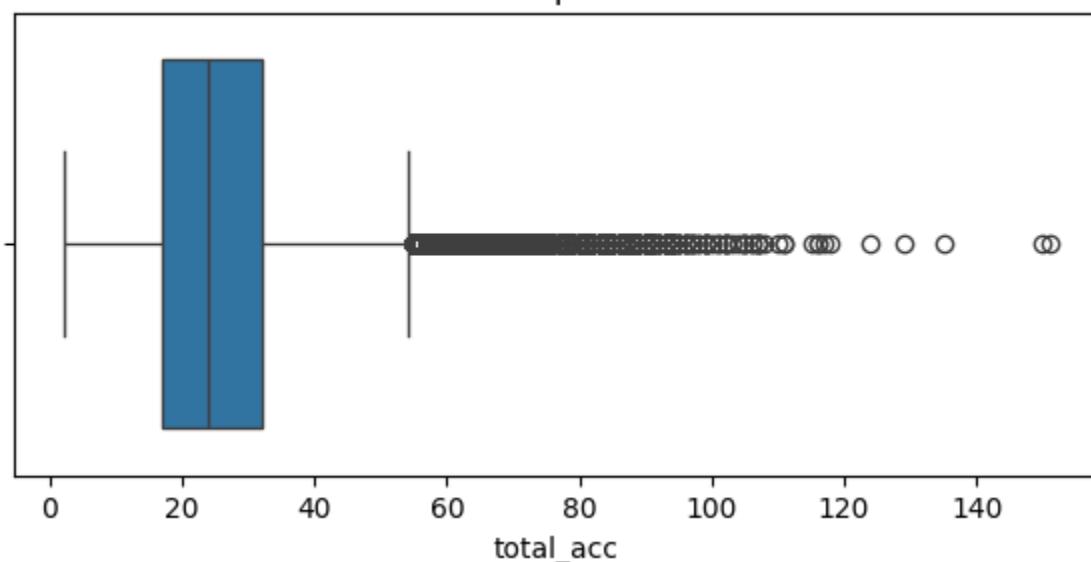
Boxplot



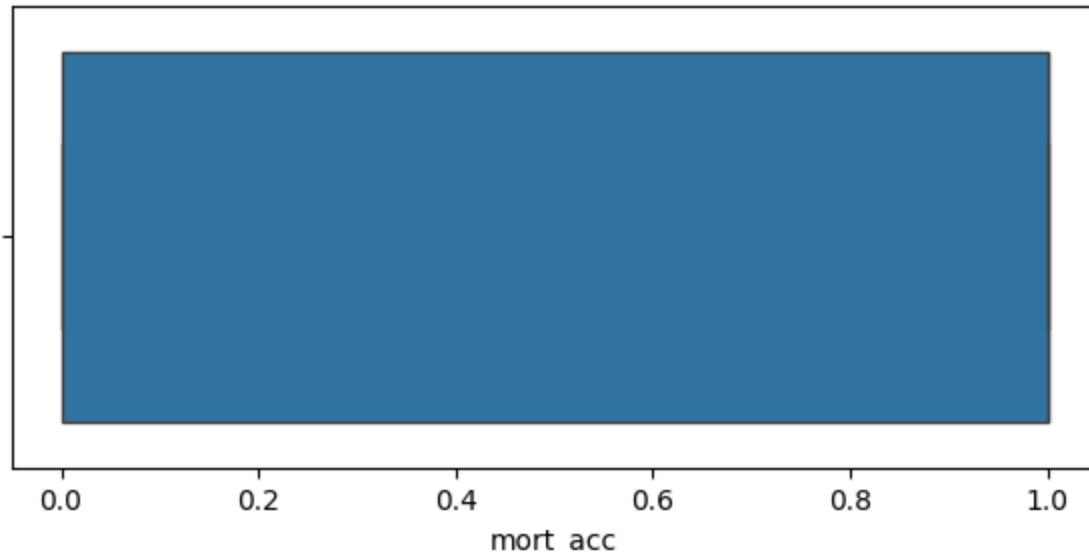
Boxplot



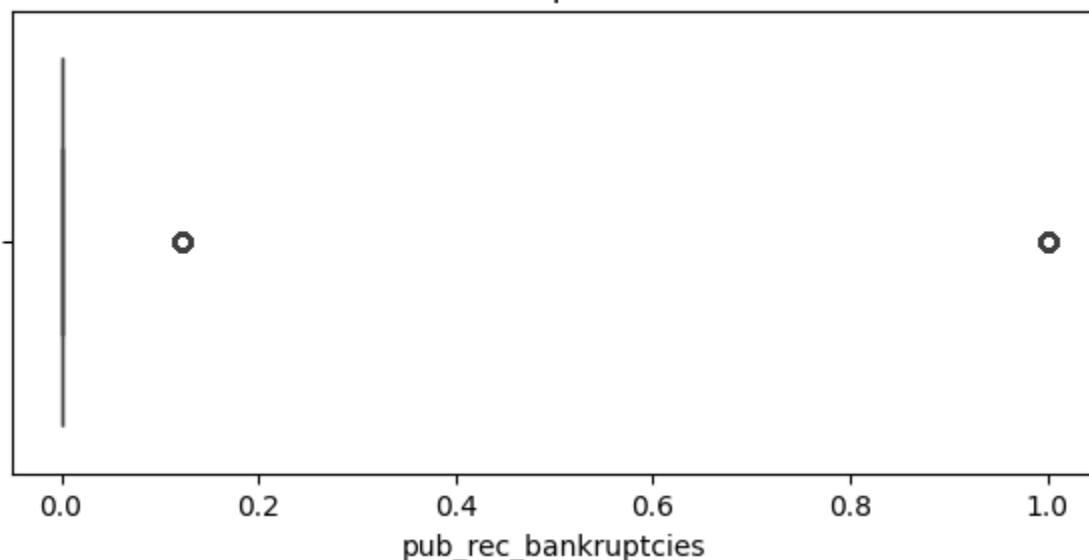
Boxplot



Boxplot



Boxplot



```
In [56]: for col in num_cols:  
    mean=df[col].mean()  
    std=df[col].std()  
  
    upper_limit=mean+3*std  
    lower_limit=mean-3*std  
  
    df=df[(df[col]<upper_limit) & (df[col]>lower_limit)]
```

## One Hot Encoding

```
In [57]: dummies=['purpose', 'zip_code', 'grade', 'verification_status', 'application_t  
df=pd.get_dummies(df,columns=dummies,drop_first=True)
```

```

pd.set_option('display.max_columns',None)
pd.set_option('display.max_rows',None)

df.head()

```

Out[57]:

	loan_amnt	term	int_rate	installment	annual_inc	loan_status	dti	open_ac	pub_ac
<b>0</b>	10000.0	36	11.44	329.48	117000.0		0	26.24	1
<b>1</b>	8000.0	36	11.99	265.68	65000.0		0	22.05	1
<b>2</b>	15600.0	36	10.49	506.97	43057.0		0	12.79	1
<b>3</b>	7200.0	36	6.49	220.65	54000.0		0	2.60	1
<b>4</b>	24375.0	60	17.27	609.33	55000.0		1	33.95	1

## Train-Test Split

In [58]:

```

from sklearn.model_selection import train_test_split
X=df.drop('loan_status',axis=1)
y=df['loan_status']
X_train , X_test , y_train , y_test = train_test_split(X,y,random_state=3,test_size=0.2)
print(X_train.shape)
print(X_test.shape)

```

(299649, 52)  
(74913, 52)

## MinMaxScaler

In [59]:

```

from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train = pd.DataFrame(scaler.fit_transform(X_train), columns=X_train.columns)
X_test = pd.DataFrame(scaler.transform(X_test), columns=X_test.columns)

```

In [60]:

```
X_train.head()
```

Out[60]:

	loan_amnt	term	int_rate	installment	annual_inc	dti	open_acc	pub_ac
<b>0</b>	0.429230	0.0	0.518228	0.497159	0.233871	0.284874	0.461538	
<b>1</b>	0.266623	0.0	0.461006	0.303769	0.229032	0.248459	0.346154	
<b>2</b>	0.039500	0.0	0.359483	0.044141	0.125000	0.267927	0.269231	
<b>3</b>	0.232390	0.0	0.106138	0.235643	0.147056	0.263725	0.230769	
<b>4</b>	0.131007	0.0	0.059991	0.130509	0.177419	0.135574	0.269231	

# Oversampling with SMOTE

```
In [61]: from imblearn.over_sampling import SMOTE
sm=SMOTE(random_state=42)
X_train_res, y_train_res = sm.fit_resample(X_train,y_train.ravel())

print(f"Before OverSampling, count of label 1: {sum(y_train == 1)}")
print(f"Before OverSampling, count of label 0: {sum(y_train == 0)}")
print(f"After OverSampling, count of label 1: {sum(y_train_res == 1)}")
print(f"After OverSampling, count of label 0: {sum(y_train_res == 0)}")
```

```
Before OverSampling, count of label 1: 58851
Before OverSampling, count of label 0: 240798
After OverSampling, count of label 1: 240798
After OverSampling, count of label 0: 240798
```

# Logistic Regression

```
In [62]: from sklearn.linear_model import LogisticRegression

model = LogisticRegression()
model.fit(X_train_res, y_train_res)
train_preds = model.predict(X_train)
test_preds = model.predict(X_test)
```

```
In [66]: from sklearn.metrics import (accuracy_score, confusion_matrix, roc_curve, auc,
                                      f1_score, recall_score,
                                      precision_score, precision_recall_curve,
                                      average_precision_score, classification_report)

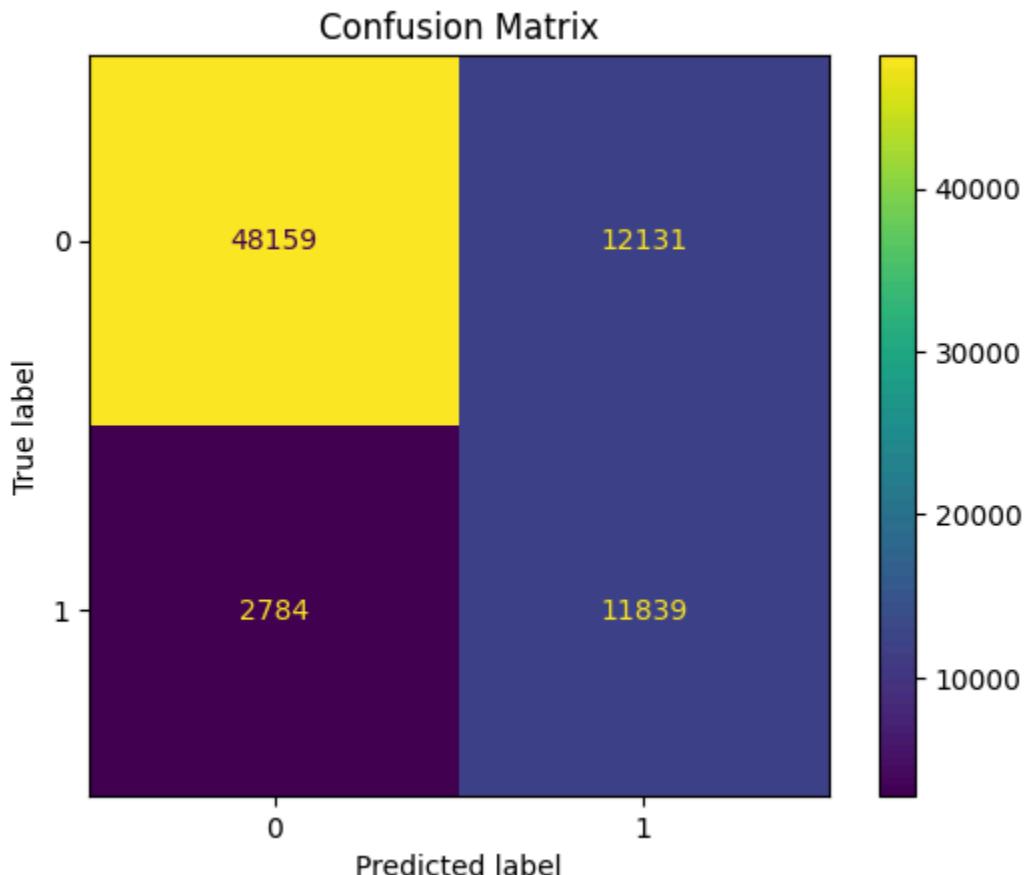
#Model Evaluation
print('Train Accuracy :', round(model.score(X_train, y_train), 2))
print('Train F1 Score:', round(f1_score(y_train, train_preds), 2))
print('Train Recall Score:', round(recall_score(y_train, train_preds), 2))
print('Train Precision Score:', round(precision_score(y_train, train_preds), 2))

print('\nTest Accuracy :', round(model.score(X_test,y_test), 2))
print('Test F1 Score:', round(f1_score(y_test,test_preds), 2))
print('Test Recall Score:', round(recall_score(y_test,test_preds), 2))
print('Test Precision Score:', round(precision_score(y_test,test_preds), 2))
```

```
Train Accuracy : 0.8
Train F1 Score: 0.61
Train Recall Score: 0.81
Train Precision Score: 0.49
```

```
Test Accuracy : 0.8
Test F1 Score: 0.61
Test Recall Score: 0.81
Test Precision Score: 0.49
```

```
In [67]: # Confusion Matrix
cm = confusion_matrix(y_test, test_preds)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```



## Classification Report

```
In [68]: print(classification_report(y_test, test_preds))
```

	precision	recall	f1-score	support
0	0.95	0.80	0.87	60290
1	0.49	0.81	0.61	14623
accuracy			0.80	74913
macro avg	0.72	0.80	0.74	74913
weighted avg	0.86	0.80	0.82	74913

- We notice that the recall score is notably high, indicating our model can identify 80% of actual defaulters. However, the precision for the

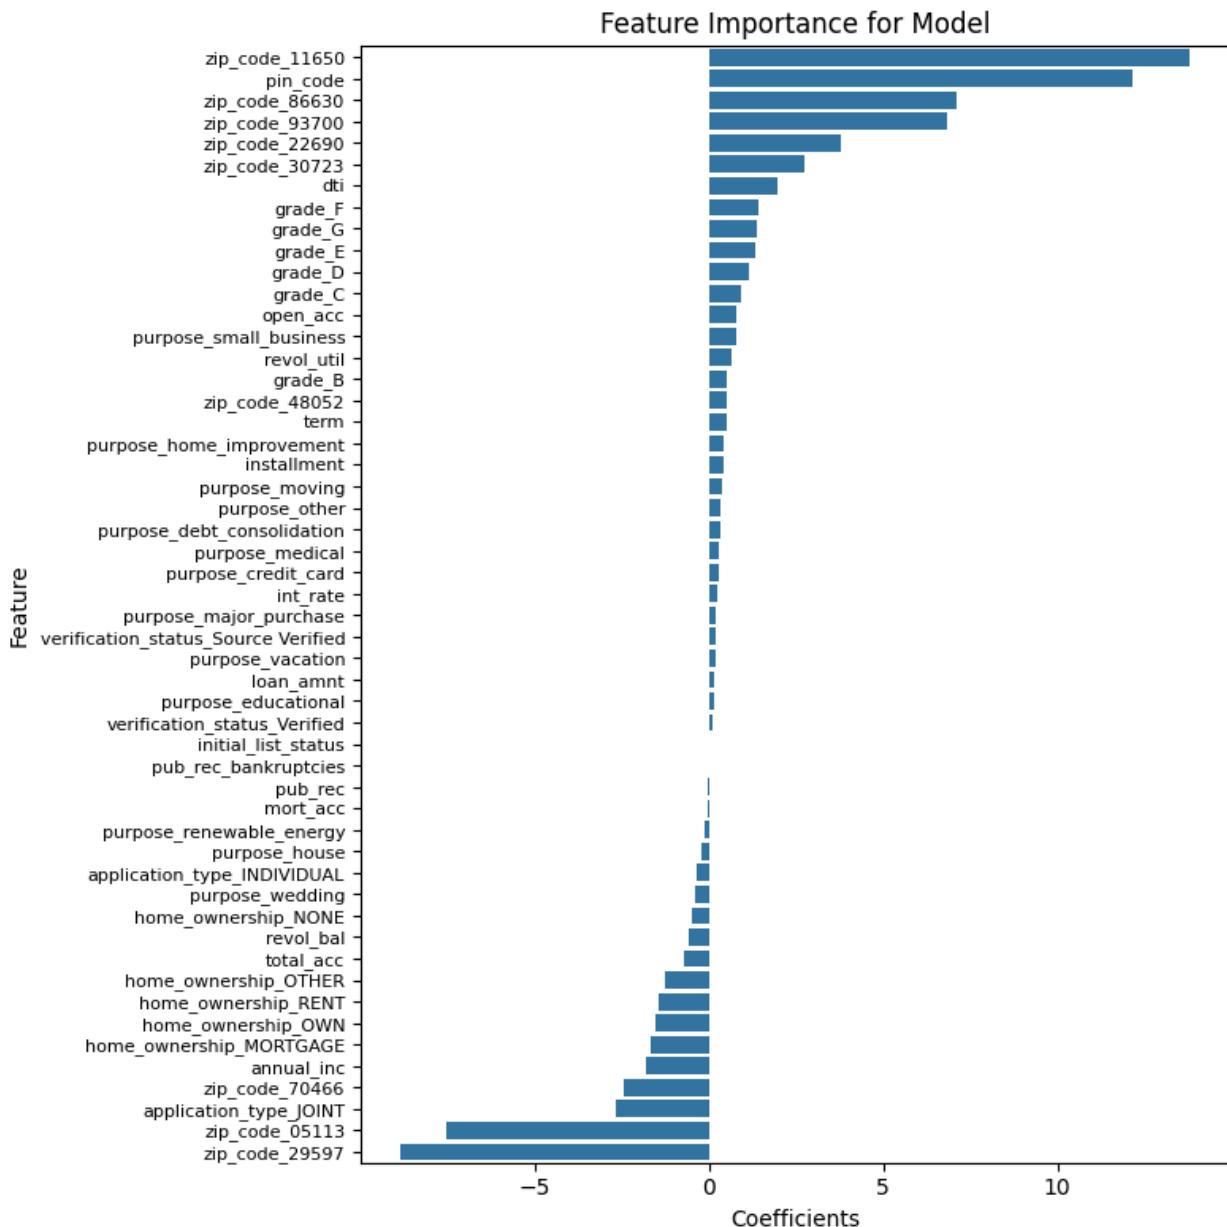
positive class is low; only 50% of the predicted defaulters are actual defaulters.

- While this model effectively identifies a significant portion of defaulters, it risks denying loans to deserving customers due to the high rate of false positives.
- Furthermore, the low precision contributes to a decrease in the F1 score to 60%, despite the accuracy being 80%.

## Feature Importance

```
In [69]: feature_imp = pd.DataFrame({'Columns':X_train.columns,
                                         'Coefficients':model.coef_[0]}).round(2).sort_values

plt.figure(figsize=(8,8))
sns.barplot(y = feature_imp['Columns'],
             x = feature_imp['Coefficients'])
plt.title("Feature Importance for Model")
plt.yticks(fontsize=8)
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



## ROC Curve & AUC

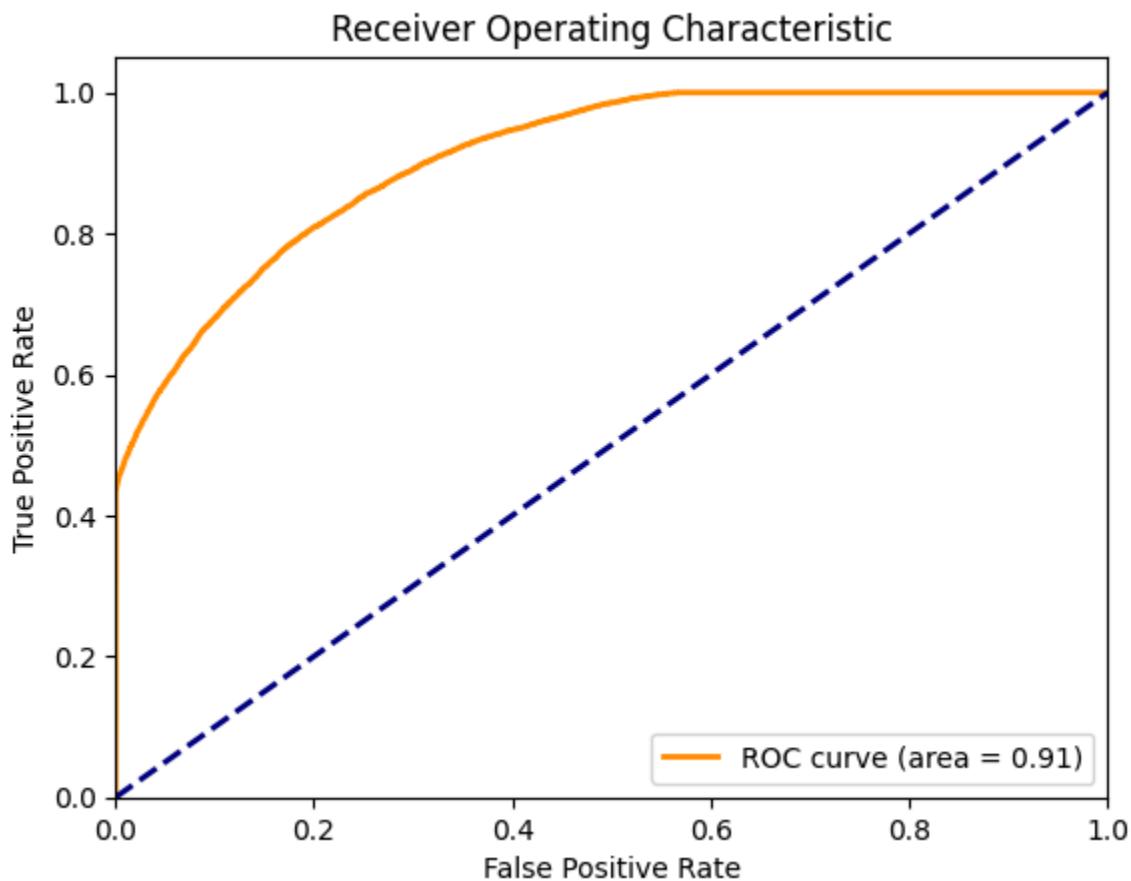
```
In [70]: # Predict probabilities for the test set
probs = model.predict_proba(X_test)[:,1]

# Compute the false positive rate, true positive rate, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, probs)

# Compute the area under the ROC curve
roc_auc = auc(fpr, tpr)

# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
```

```
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



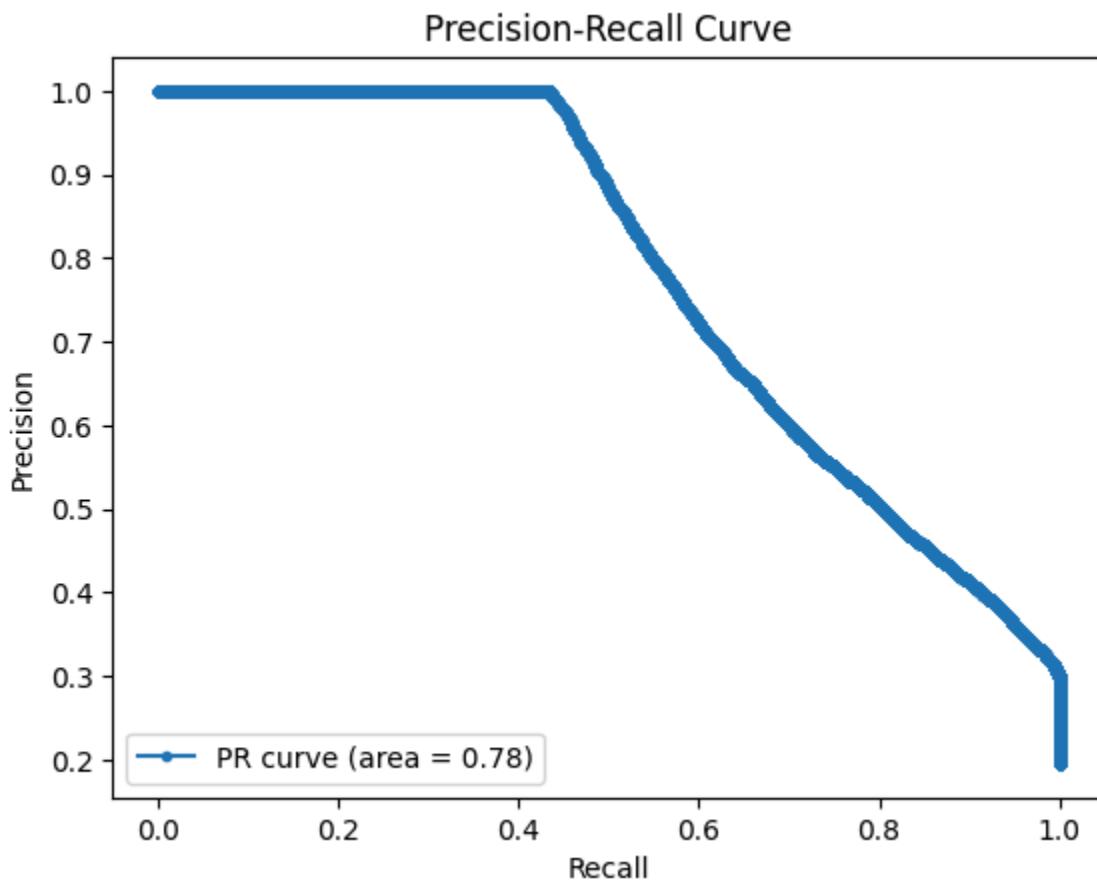
- An AUC of 0.91 indicates that the model effectively distinguishes between the positive and negative classes.
- However, it's not an ideal metric for imbalanced target variables since it can be high even when the classifier performs poorly on the minority class.
- This discrepancy occurs when the classifier excels at classifying instances from the majority class, which are more prevalent in the dataset. Consequently, the AUC might appear high, masking the model's inability to accurately identify instances from the minority class.

## Precision Recall Curve

```
In [71]: precision, recall, thr = precision_recall_curve(y_test, probs)

# Area under Precision Recall Curve
apc = average_precision_score(y_test, probs)

# Plot the precision-recall curve
plt.plot(recall, precision, marker='.', label='PR curve (area = %0.2f)' % apc)
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc="lower left")
plt.show()
```



The area under the precision-recall curve (AUC-PR) is not as high as desired. While it exceeds the benchmark of 0.5 for a random model, indicating some level of effectiveness.

## Conclusion

**Q1.** How can we make sure that our model can detect real defaulters and there are

less false positives? This is important as we can lose out on an opportunity to finance more individuals and earn interest on it.

**Ans:** The precision score serves as an indicator of Type I error. By increasing the precision score of the model, we can minimize false positives. This ensures that the company avoids erroneously denying loans to deserving individuals, thus maximizing the opportunity to finance worthy applicants.

**Q2.** Since NPA (non-performing asset) is a real problem in this industry, it's important we play safe and shouldn't disburse loans to anyone.

**Ans:** The recall score serves as an indicator of how effectively the model identifies actual defaulters. By increasing the recall score, we can minimize false negatives (Type II error), thereby ensuring that loans are not disbursed to defaulters, thus enhancing the model's ability to identify risky applicants.

## Insights

- 80% belongs to the class 0 : which is loan fully paid.
- 20% belongs to the class 1 : which were charged off.
- Loan Amount distribution / media is slightly higher for Charged\_off loanStatus.
- the probability of defaulters is higher in the small\_business owner borrowers.
- Total credit revolving balance is almost same for both borrowers who had fully paid loan and declared defaulter
- Probability of Charged\_off status is higher in case of 60 month term.
- Zip codes such as 11650, 86630 and 93700 have 100% defaulters
- It can be observed that the mean loan\_amnt, int\_rate, dti, open\_acc and revol\_util are higher for defaulters.
- The % of defaulters is much higher for longer (60-month) term.
- A Logistic Regression model performed well, rendering accuracy of 80%.
- We can remove initial\_list\_status and state as they have no impact on

### loan\_status

- The model had a precision score of 95%, recall score of 80%, and f1 score of 87% on the negative class.
- The model had a precision score of 49%, recall score of 81%, and f1 score of 61% on the positive class.
- The features "grade" and "sub-grade" have the most significant impact on the loan\_status, with higher grades typically associated with a higher likelihood of default. In particular, loans assigned the highest grade tend to have the highest proportion of defaulters.

## Recommendations

- Since NPA is a real problem in the industry , Company should more investigate and check for the proof of assets. Since it was observed in probability plot, verified borrowers had higher probability of defaulters than non-varified.
- Prioritize 'A' grade applicants and shorter-term loans for lower default risk.
- Balancing risk of increasing NPAs by disbursing loans to defaulters with the opportunity to earn interest by disbursing loans to as many worthy customers as possible is to maximize the F1 score along with the area under the Precision-Recall Curve.